

AWS Autoscaling Budget Estimate:

Depending on the number of requests received and the utilization percentage of the CPU of a single instance (server), if the CPU is being utilized >85%, AWS auto-scales and clones and starts a new instance immediately.

The instance type we are using here is t4g.nano which costs \$0.0028 per hour per instance. T4g.nano instance has 2 vCPUs and 1GB RAM, enough to serve a very huge number of requests concurrently, without a lag. Even if there are >8k requests being sent to the servers, it would basically cost $2 \times 0.0028 \times 24$ if the two instances having >85% CPU utilization are being run simultaneously. Running a single t4g.nano instance for a whole month would cost us 2.45USD.

If we run 2 servers simultaneously for a week (considering the FFCS days) it would cost us 1.8816USD, which is 69.94 INR, which is nothing.

There will not be any downtimes, no rate limits and absolutely cost-efficient.